



General

FA2025 • 2025-10-09

Reverse Engineering II

Cameron and Satvik

Announcements

- Pizza and Game Night tomorrow night!
 - This is taking the place of Amateurs CTF, which was moved to November.
 - Details will be sent in an announcement on Discord



ctf.sigpwny.com


sigpwny{rev_your_engines}

WHAT MY CODE SAYS

```
float get_biggest_number(float a, float b){  
    bool is_a_biggest;  
    bool is_b_biggest;  
    if (a > b){  
        is_a_biggest = true;  
    }  
    else {  
        is_a_biggest = false;  
    }  
    if (b > a){  
        is_b_biggest = true;  
    }  
    else {  
        is_b_biggest = false;  
    }  
    if (is_a_biggest == true){  
        return a;  
    }  
    if (is_b_biggest == true){  
        return b;  
    }  
}
```


WHAT COMPILER THINKS:

```
1 get_biggest_number(float, float):  
2     maxss    xmm0, xmm1  
3     ret
```



GCC-03

"Sometimes my genius is... it's almost frightening"



Setup

- If you haven't installed Ghidra yet, start downloading it through the slides here: sigpwny.com/rev_setup



Recap: Reverse Engineering

- Reverse Engineering: Figure out how a program works
 - more broadly: get useful information out of a program
- Why reverse engineering?
 - Solve reverse engineering CTF challenges and get flags
 - Find vulnerabilities in software
 - Makes you a better programmer
 - And more
- Two major (non-exclusive) techniques
 - Static analysis (today: **Ghidra**)
 - Dynamic analysis (today: **GDB**)



Recap: Assembly

Last meeting: sigpwny.com/meetings/general/2025-10-05/



What is Assembly?

- A human-readable abstraction over CPU machine codes

01001000000001011101111011000000011011100010011

48 05 DE C0 37 13

add rax, 0x1337c0de



What is Assembly?

```
int method(int a){  
    int b = 6;  
    char c = 'c';  
    return a+b;  
}
```

method:

```
push    rbp  
mov     rbp, rsp  
mov     DWORD PTR [rbp-20], edi  
mov     DWORD PTR [rbp-4], 6  
mov     BYTE PTR [rbp-5], 99  
mov     edx, DWORD PTR [rbp-20]  
mov     eax, DWORD PTR [rbp-4]  
add     eax, edx  
pop     rbp  
ret
```



Basic CPU Structures

Instruction Memory

```
[0x00401000]
    ;-- section..text:
    ;-- segment.LOAD1:
entry0 ();
push    rsp
pop     rsi
xor     dl, 0x60
syscall
ret
```

Registers

```
*RAX 0x3e8
*RBX 0x401300 (__libc_csu_init) ←
*RCX 0x7ffff7ea311b (getegid+11) ←
RDX 0x0
*RDI 0x7ffff7fad7e0 (_IO_stdfile_1
RSI 0x0
R8 0x0
*R9 0x7ffff7fe0d60 (_dl_fini) ←
*R10 0x400502 ← 0x64696765746567
*R11 0x202
*R12 0x401110 (_start) ← endbr64
*R13 0x7ffffffffffddc0 ← 0x1
R14 0x0
R15 0x0
*RBP 0x7ffffffffffdcd0 ← 0x0
*RSP 0x7ffffffffffdcb0 ← 0x0
*RIP 0x401220 (main+42) ← mov
```

Stack

```
0x7ffffffffffdcb0 ← 0x0
0x7ffffffffffdcb8 → 0x401110 (_start)
0x7ffffffffffdcc0 → 0x7ffffffffffddc0
0x7ffffffffffdcc8 ← 0x0
0x7ffffffffffdcd0 ← 0x0
0x7ffffffffffdcd8 → 0x7ffff7de3083
```



What is this meeting about?

- Reverse engineering **binaries**
 - Compiled executables
 - All source information is usually (but not always) stripped
- What do we have to work with?
 - Machine code
 - Sometimes, some symbol names (like function names)
 - At minimum, only what the OS needs to execute the program



Running example: debugger

```
→ rev ./debugger sigpwny{test_flag}  
That flag is incorrect.  
→ rev █
```

- Challenge might feel completely opaque right now
- But we will be able to solve it by the end of the meeting
- Follow along!



The ELF Format

- What kind of file is debugger?
 - The more information you have about the program you are reversing, the easier it is
- Use Unix “file” utility

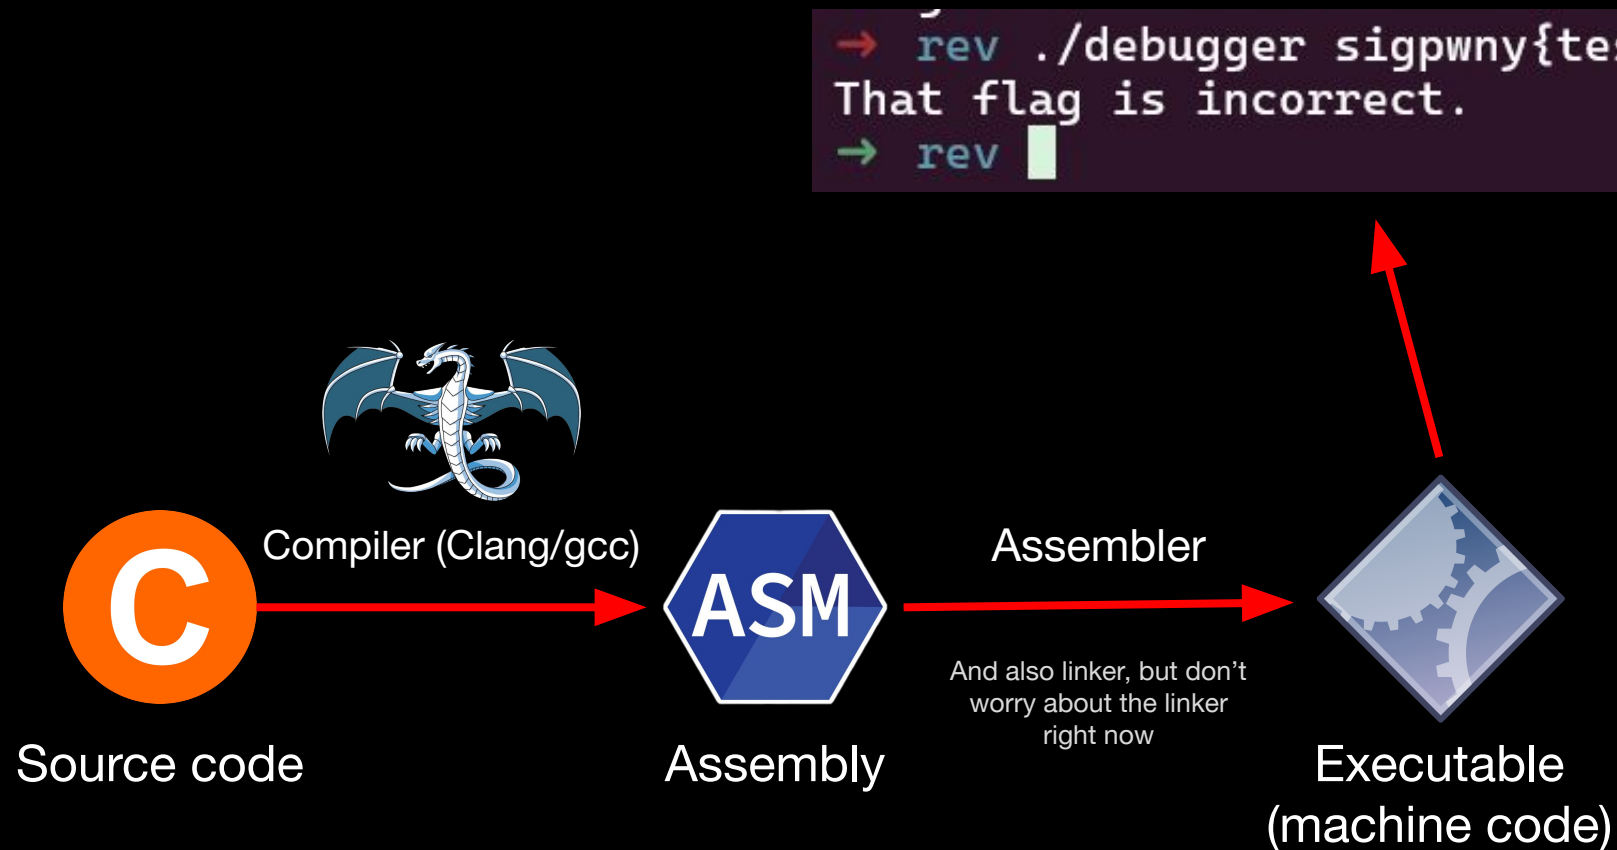
```
→ rev file debugger
debugger: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically
linked, interpreter /lib64/ld-linux-x86-64.so.2, BuildID[sha1]=7b85de3d4b
fac967613aa60d4d1540f90e5d8676, for GNU/Linux 3.2.0, not stripped
```

- ELF: Executable and Linkable Format
 - File format for **executables**, libraries, object files
 - Contains program code and data, plus metadata needed to execute program
 - Can also contain symbols (“not stripped”)
 - More info:
<https://github.com/corkami/pics/blob/28cb0226093ed57b348723bc473cea0162dad366/binary/elf101/elf101.pdf>
 - Useful tool: readelf

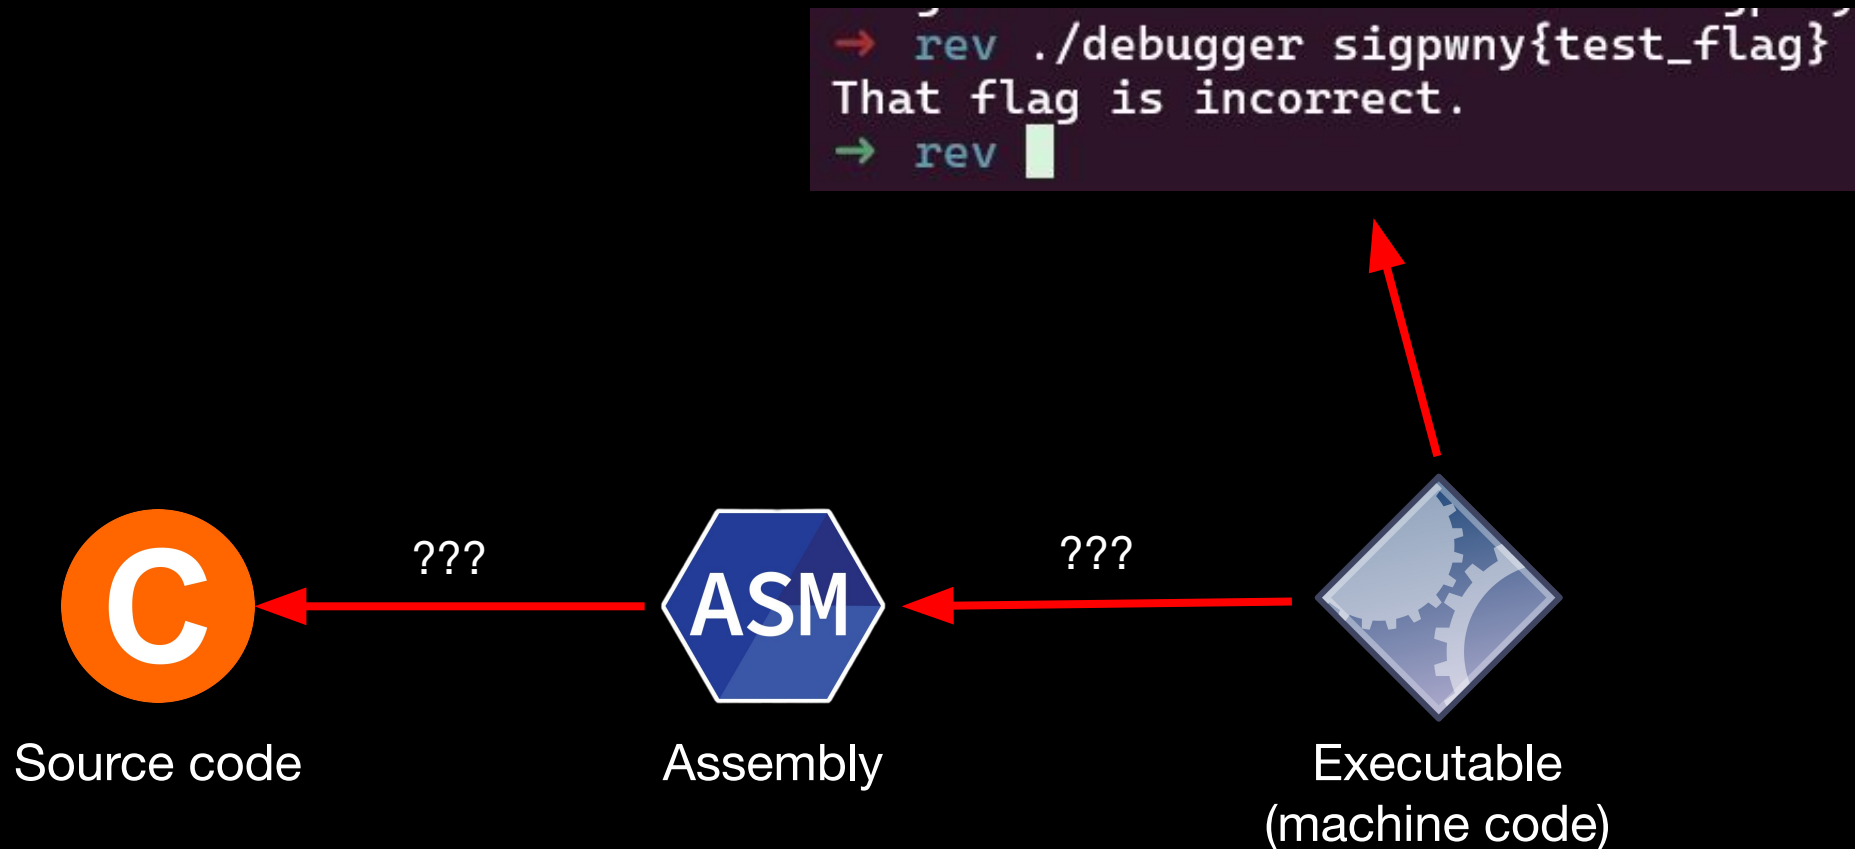


Compilation

Or, how does source code become an executable




Can we go the other way?



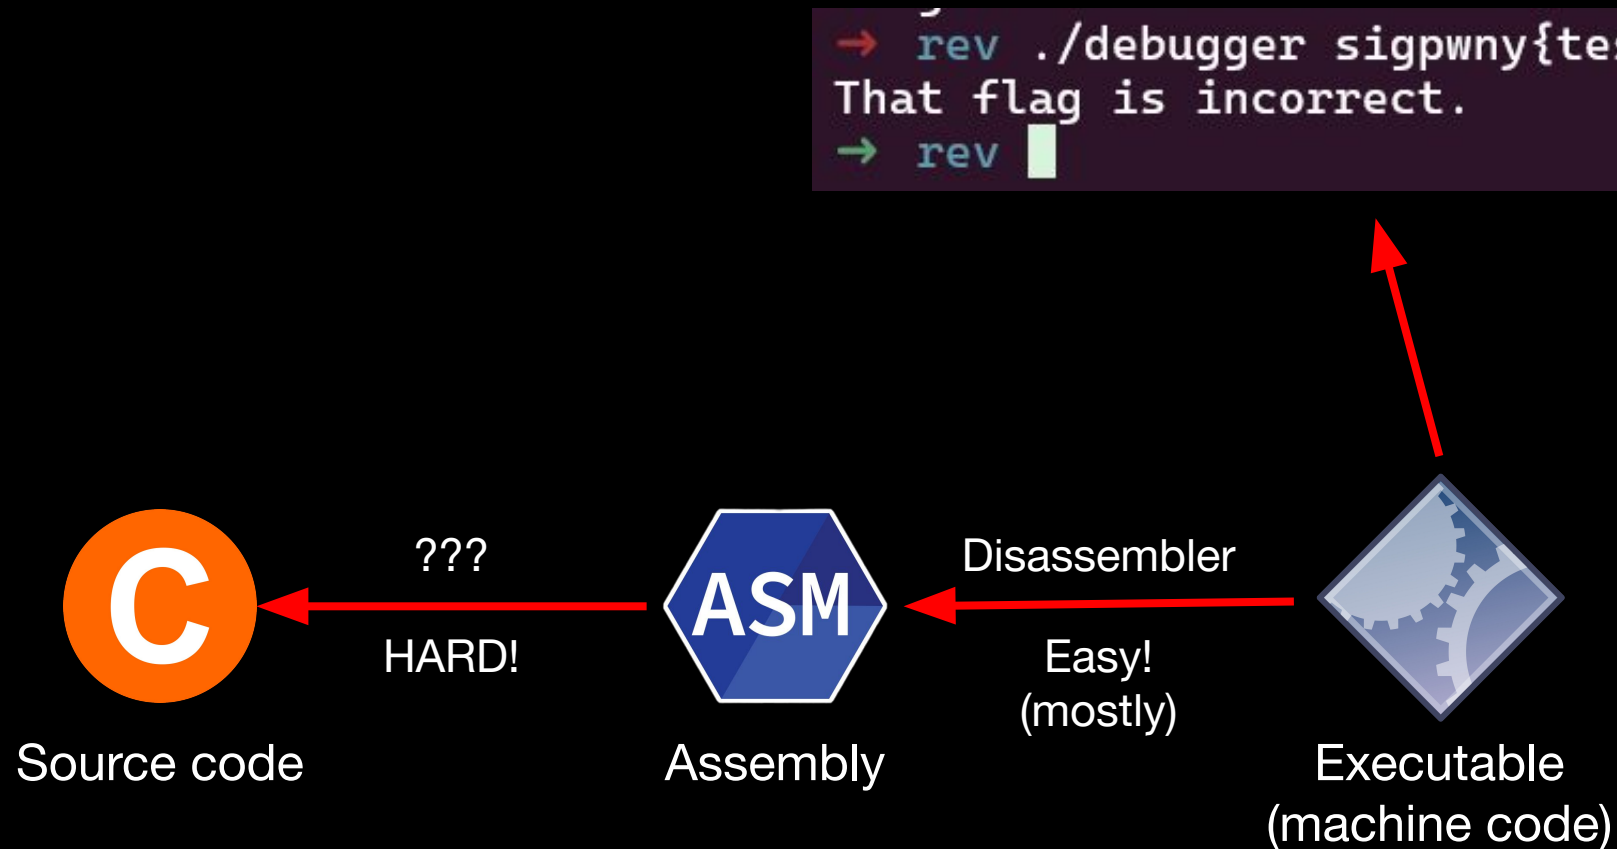
Can we go the other way?

```
pwndbg> disass main
Dump of assembler code for function main:
0x0000000000401150 <+0>:    push    rbp
0x0000000000401151 <+1>:    mov     rbp, rsp
0x0000000000401154 <+4>:    sub     rsp, 0x40
0x0000000000401158 <+8>:    mov     DWORD PTR [rbp-0x4], 0x0
0x000000000040115f <+15>:   mov     DWORD PTR [rbp-0x8], edi
0x0000000000401162 <+18>:   mov     QWORD PTR [rbp-0x10], rsi
0x0000000000401166 <+22>:   cmp     DWORD PTR [rbp-0x8], 0x2
0x000000000040116a <+26>:   jge     0x40118b <main+59>
0x0000000000401170 <+32>:   movabs  rdi, 0x402004
0x000000000040117a <+42>:   call    0x401040 <puts@plt>
```

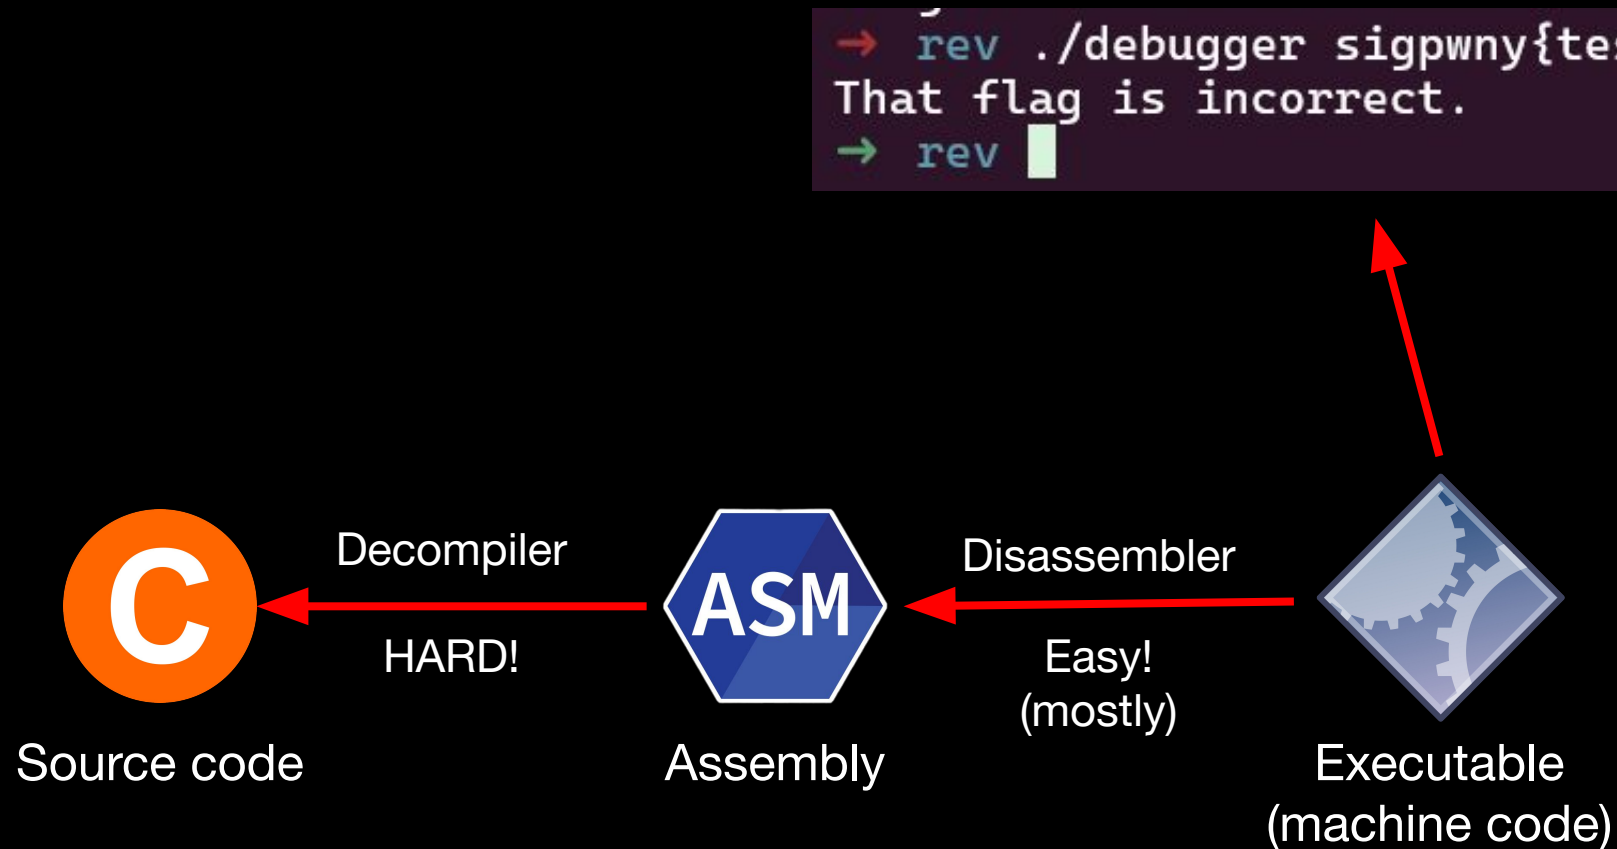
```
→ rev ./debugger sigpwny{test_flag}
That flag is incorrect.
→ rev
```



Can we go the other way?



Can we go the other way?



Decompilation



We can go from C code to assembly...

```
1 int some_mathz() {  
2     int res = 0;  
3     for (int i = 9; i > 1; i++) {  
4         res -= i;  
5     }  
6 }
```

some_mathz():

```
    push    rbp  
    mov     rbp, rsp
```

```
    mov     DWORD PTR [rbp-4], 0
```

```
    mov     DWORD PTR [rbp-8], 9
```

```
    jmp     .L2
```

.L3:

```
    mov     eax, DWORD PTR [rbp-8]
```

```
    sub     DWORD PTR [rbp-4], eax
```

```
    add     DWORD PTR [rbp-8], 1
```

.L2:

```
    cmp     DWORD PTR [rbp-8], 1
```

```
    jg      .L3
```

```
    ud2
```



Now go from assembly to C code 🤖

Challenge: What does this do?

```
1  add(unsigned int):  
2      test    edi, edi  
3      je      .L4  
4      mov     eax, 1  
5      mov     edx, 0  
6  .L3:  
7      add     edx, eax  
8      add     eax, 1  
9      cmp     edi, eax  
10     jnb     .L3  
11  .L2:  
12     mov     eax, edx  
13     ret  
14  .L4:  
15     mov     edx, edi  
16     jmp     .L2
```



Now go from assembly to C code 🐉

Challenge: What does this do?

```
1  add(unsigned int):
2      test    edi, edi
3      je      .L4
4      mov     eax, 1
5      mov     edx, 0
6  .L3:
7      add     edx, eax
8      add     eax, 1
9      cmp     edi, eax
10     jnb     .L3
11  .L2:
12     mov     eax, edx
13     ret
14  .L4:
15     mov     edx, edi
16     jmp     .L2
```

```
unsigned add(unsigned n) {
    // Compute 1 + 2 + ... + n
    unsigned result = 0;
    for (unsigned i = 1; i <= n; i++) {
        result += i;
    }
    return result;
}
```



Ghidra to the rescue!

- Open source disassembler/decompiler/"reverse engineering framework"
 - **Disassembler**: binary machine code to assembly
 - **Decompiler**: assembly to pseudo-C
 - Reverse engineering framework: control flow graph recovery, cross-references, binary similarity/diffing, and more!
- Written by the NSA 🤖



Ghidra caveats

```
unsigned add(unsigned n) {  
    // Compute 1 + 2 + ... + n  
    unsigned result = 0;  
    for (unsigned i = 1; i <= n; i++) {  
        result += i;  
    }  
    return result;  
}
```

Decompilation not always the same! Many ways to write equivalent code

```
uint add(uint n)  
{  
    uint i;  
    uint result;  
  
    result = n;  
    if (n != 0) {  
        i = 1;  
        result = 0;  
        do {  
            result = result + i;  
            i = i + 1;  
        } while (i <= n);  
    }  
    return result;  
}
```



Ghidra caveats

- Ghidra output is not meant to be recompileable
 - It's meant to be human-readable
- Decompilation is a best guess
 - But not all information (e.g. types) is always recovered

```
1
2 undefined4 main(int argc, char **argv)
3
4 {
5     int iVar1;
6     size_t sVar2;
7     uint local_44;
8     undefined8 local_40;
9     undefined8 local_38;
10    undefined8 local_30;
11    undefined4 local_28;
12    undefined local_24;
13    char **local_18;
14    int local_10;
15    undefined4 local_c;
16
```



Common Optimizations

Loading an array with bytes

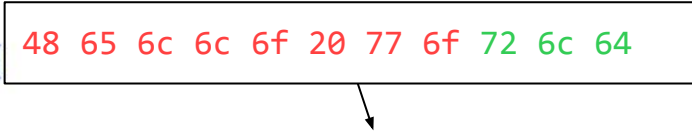
- Loading first 8 bytes simultaneously into stack (in one instruction)

```
#include <stdio.h>

int main() { 48 65 6c 6c 6f 20 77 6f 72 6c 64

    char string[] = "Hello world";
    printf("%s",string);

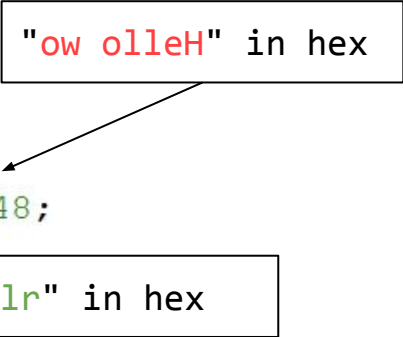
    return 0;
}
```



Challenge: why is the text of the decoded number backwards?

```
int __cdecl main(int _Argc, char **_Argv, char **_Env)
{
    undefined8 local_14;
    undefined4 local_c;

    __main();
    local_14 = 0x6f77206f6c6c6548;
    local_c = 0x646c72;
    printf("%s",&local_14);
    return 0;
}
```



Common Optimizations (Cont.)

Modulo replaced with mask

- % 4 replaced with & 0b11 (Taking the last two bits of unsigned int)

```
#include <stdio.h>

int main() {

    unsigned int A,B;
    scanf("%u",&A);
    B = A % 4;
    printf("%u",B);

    return 0;
}
```

```
int __cdecl main(int _Argc,char **_Argv,char **_Env)

{
    uint A;
    uint B;

    __main();
    scanf("%u",&A);
    B = A & 0b00000011;
    printf("%u", (ulonglong)B);
    return 0;
}
```

Ghidra Cheat Sheet

- Get started:
 - View all functions in list on left side of screen inside “Symbol Tree”. Double click **main** to decompile main
- Decompiler:
 - Middle click a variable to highlight all instances in decompilation
 - Type “L” to rename variable (after clicking on it)
 - “Ctrl+L” to retype a variable (type your type in the box)
 - Type “;” to add an inline comment on the decompilation and assembly
 - Alt+Left Arrow to navigate back to previous function
- General:
 - Double click an XREF to navigate there
 - Search -> For Strings -> Search to find all strings (and XREFs)
 - Choose Window -> Function Graph for a graph view of disassembly



GDB (Dynamic Analysis)

- Able to inspect a program's variables & state as it runs
- Set breakpoints, step through, try various inputs
- Debugging analogy: print statements after running



Dynamic Analysis with GDB

- Run program, with the ability to pause and resume execution
- View registers, stack, heap
- Steep learning curve
- `chmod +x ./chal` to make executable

```
B+ 0x55555555129 <add>          endbr64
0x5555555512d <add+4>          test    %edi,%edi
0x5555555512f <add+6>          je      0x55555555147 <add+30>
0x55555555131 <add+8>          mov     $0x1,%eax
0x55555555136 <add+13>         mov     $0x0,%edx
0x5555555513b <add+18>         add     %eax,%edx
0x5555555513d <add+20>         add     $0x1,%eax
> 0x55555555140 <add+23>        cmp     %eax,%edi
0x55555555142 <add+25>        jae     0x5555555513b <add+18>
0x55555555144 <add+27>        mov     %edx,%eax
0x55555555146 <add+29>        retq
0x55555555147 <add+30>        mov     %edi,%edx
0x55555555149 <add+32>        jmp     0x55555555144 <add+27>
0x5555555514b <main>          endbr64
0x5555555514f <main+4>         callq   0x55555555129 <add>
0x55555555154 <main+9>         retq
0x55555555155                nopw    %cs:0x0(%rax,%rax,1)
0x5555555515f                nop
0x55555555160 <__libc_csu_init>        endbr64
0x55555555164 <__libc_csu_init+4> push    %r15

native process 219424 In: add
rax      0x4      4
rbx      0x55555555160  93824992235872
rcx      0x55555555160  93824992235872
rdx      0x6      6
rsi      0x7fffffffdd58  140737488346456
--Type <RET> for more, q to quit, c to continue without paging--
```

GDB Cheat Sheet

[gdb](#)

[pwndbg](#)

- `b main` - Set a breakpoint on the main function
 - `b *main+10` - Set a breakpoint a couple instructions into main
- `r` - run
 - `r arg1 arg2` - Run program with arg1 and arg2 as command line arguments. Same as `./prog arg1 arg2`
 - `r < myfile` - Run program and supply contents of myfile.txt to stdin
- `c` - continue
- `si` - step instruction (steps into function calls)
- `ni` - next instruction (steps over function calls) (`finish` to return to caller function)
- `x/32xb 0x5555555551b8` - Display 32 hex bytes at address 0x5555555551b8
 - `x/4xg addr` - Display 4 hex “giants” (8 byte numbers) at addr
 - `x/16i $pc` - Display next 16 instructions at \$rip
 - `x/s addr` - Display a string at address
 - `x/4gx {void*}$rcx` - Dereference pointer at \$rcx, display 4 QWORDS
 - `p/d {int*}{int*}$rcx` - Dereference pointer to pointer at \$rcx as decimal
- `info registers` - Display registers (shorthand: `i r`)
- `x86 Linux calling convention`* (“System V ABI”): RDI, RSI, RDX, RCX, R8, R9

*syscall calling convention is RDI, RSI, RDX, **R10**, R8, R9



pwndbg

git clone

<https://github.com/pwndbg/pwndbg>

cd pwndbg

./setup.sh

Breakpoint 1, 0x0000000000401150 in main ()
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA

[REGISTERS]

```
RAX 0x401150 (main) ← push rbp
RBX 0x0
RCX 0x401290 (__libc_csu_init) ← endbr64
RDX 0x7fffffffef1a8 → 0x7fffffffef49a ← 'DBUS_SESSION_BUS_ADDRESS=unix:path=/run/user/1000/bus'
RDI 0x1
RSI 0x7fffffffef198 → 0x7fffffffef47d ← '/home/richyliu/temp/debugger'
R8 0x7ffff7f90f10 (initial+16) ← 0x4
R9 0x7ffff7fc9040 (_dl_fini) ← endbr64
R10 0x7ffff7fc3908 ← 0xd00120000000e
R11 0x7ffff7fde680 (_dl_audit_preinit) ← endbr64
R12 0x7fffffffef198 → 0x7fffffffef47d ← '/home/richyliu/temp/debugger'
R13 0x401150 (main) ← push rbp
R14 0x0
R15 0x7ffff7ffd040 (_rtld_global) → 0x7ffff7ffe2e0 ← 0x0
RBP 0x1
RSP 0x7fffffffef088 → 0x7ffff7d9fd90 (__libc_start_call_main+128) ← mov edi, eax
RIP 0x401150 (main) ← push rbp
```

[DISASM]

```
► 0x401150 <main>      push rbp
0x401151 <main+1>     mov rbp, rsp
0x401154 <main+4>     sub rsp, 0x40
0x401158 <main+8>     mov dword ptr [rbp - 4], 0
0x40115f <main+15>    mov dword ptr [rbp - 8], edi
0x401162 <main+18>    mov qword ptr [rbp - 0x10], rsi
0x401166 <main+22>    cmp dword ptr [rbp - 8], 2
0x40116a <main+26>    jge main+59          <main+59>

0x401170 <main+32>    movabs rdi, 0x402004
0x40117a <main+42>    call puts@plt        <puts@plt>

0x40117f <main+47>    mov dword ptr [rbp - 4], 1
```

[STACK]

```
00:0000 | rsp 0x7fffffffef088 → 0x7ffff7d9fd90 (__libc_start_call_main+128) ← mov edi, eax
01:0008 |      0x7fffffffef090 ← 0x0
02:0010 |      0x7fffffffef098 → 0x401150 (main) ← push rbp
03:0018 |      0x7fffffffef0a0 ← 0x100000000
04:0020 |      0x7fffffffef0a8 → 0x7fffffffef198 → 0x7fffffffef47d ← '/home/richyliu/temp/debugger'
05:0028 |      0x7fffffffef0b0 ← 0x0
06:0030 |      0x7fffffffef0b8 ← 0x8e4494d77c28027e
07:0038 |      0x7fffffffef0c0 → 0x7fffffffef198 → 0x7fffffffef47d ← '/home/richyliu/temp/debugger'
```

pwndbg> █

pwndbg cheat sheet

- `emulate #` - Emulate the next # instructions
- `stack #` - Print # values on the stack
- `vmmap` - Print memory segments (use `-x` flag to show only executable segments)
- `nearpc` - Disassemble near the PC
- `tel <ptr>` - Recursively dereferences <ptr>
- `regs` - Use instead of `info reg` (gdb's register viewing)



GDB/pwndbg for macos users

- GDB only:
 - Make sure you have docker installed and running
 - docker pull sigpwny/pwn-docker
 - One time
 - docker run -it sigpwny/pwn-docker
 - You can download the challenge files with curl -O file_url
- With pwndbg
 - <https://github.com/sigpwny/pwn-docker>
 - Follow the instructions in the README



Challenge Walkthrough

Open Ghidra!

sigpwny.com/rev_setup

Download "debugger" from <https://ctf.sigpwny.com/challenges>



Go try for yourself!

- <https://ctf.sigpwny.com>
- Start with Crackme 0
- Practice practice practice! Ask for help!



Going Further

- Side channels: e.g. instruction counting
- Symbolic/concolic execution
- Ghidra scripts
- Z3 and constraint solvers
- Emulation for dynamic analysis
- Taint analysis
- and more!
- Many of these will be covered in Rev III



Next Meetings

2025-10-10 • Tomorrow!

- Pizza and Game Night Social
- Come to Siebel CS for pizza and games!

2025-10-12 • This Sunday

- PWN I
- Learn about basic binary exploitation, such as stack overflows and overwriting return addresses!

2025-10-16 • Next Thursday

- Cryptography I
- We will cover the basics of cryptography, including ciphers, symmetric encryption, and more!



ctf.sigpwny.com

sigpwny{rev_your_engines}

Meeting content can be found at
sigpwny.com/meetings.

