



General

FA2025 • 2025-11-02

# Mayhem

Ronan Boyarski, Nikhil Date

# Finding bugs in binaries

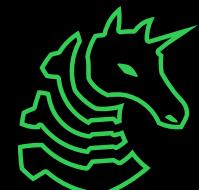
- Traditionally, you would do this by hand
- Start by using a debugger and hook the main way you write input (usually socket `recv()` function)
- Sync your debugger with your disassembler and check the control flow graph to find calls to unsafe functions
  - For example, looking for calls to `memcpy()` where you can specify the wrong input size
  - Can be subtle, like finding multiple chained `memcpy` calls where size is fixed but you can get some fixed overflow into a dangerous region
- Reliable, but slow and depends on a skilled reverse engineer
- How much of this can we automate?



# Fuzzing

- A sort of randomized testing designed to find memory bugs
- Fuzzer tries to select inputs that will explore as much of the program as possible
- The idea is that, the more of the program we explore, the higher likelihood we have of discovering a bug
- Fuzzing is very successful in practice
  - Example fuzzers include: AFL, libfuzzer, syzkaller

"More bugs than eyes. Setup Syzkaller on a junk pc tonight, by Sunday you will have unique, likely exploitable kernel bugs" -Ravi



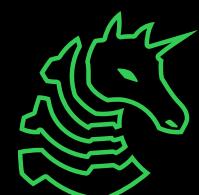
# Symbolic Execution

- Normal execution is known as "concrete" execution
- Symbolic execution: make the program inputs symbolic values
- Track symbolic expressions representing program state
- Use this to look at all possible paths
- Using symbolic execution, we can solve for whether we can mess with important bits of program state (like the instruction pointer)
- Problems:
  - Path explosion: symbolic execution does not scale well with big programs
  - Environment: how are you tracking heap allocations, syscalls, etc.
  - Reasoning about memory: how do you track memory accesses, the most important part of our manual bug finding method?



# Problem Setting

- We have a binary we want to exploit
  - Could be a really fun target, like Windows SMB server, iOS messaging app, or kernel driver
- Said binary has no debugging information, and is large and annoying
- We would like to automate finding bugs in *all* binaries of that form
  - For example, try to find exploitable bug in **every** kernel driver
- We want to automatically generate patches for all of our exploits
- We want to make sure that what the machine gives us is reliable and verifiable (no false positives > no false negatives)



# MAYHEM

- 2016 DARPA Cyber Grand Challenge winner, developed by CMU
- "Every bug reported by MAYHEM is accompanied by a working shell-spawning exploit"
- "To make exploit generation possible at the binary-level, MAYHEM addresses two major technical challenges: actively managing execution paths without exhausting memory, and reasoning about *symbolic memory indices*, where a load or a store address depends on user input. To this end, we propose two novel techniques: 1) hybrid symbolic execution for combining online and offline (concolic) execution to maximize the benefits of both techniques, and 2) index-based memory modeling."



# High-Level Idea

- What if we combine concrete execution with symbolic execution, so we can reason about real program state using symbolic execution?
- Model a formula for whether we can control the instruction pointer, put attacker supplied data in memory, and have a memory protection primitive to execute said code
  - If symbolic execution returns SAT, we can exploit the program
- How do they handle path explosion?
  - No repeated work
  - No ballooning higher than the current machine's RAM at any one time
  - Reason about symbolic memory where a load or store address depends on user input



# Hybrid Symbolic Execution

- Execution alternates between online and offline symbolic runs
- Acts like a virtual memory manager in an OS, where we swap out symbolic execution engines
  - When we are overburdened, we swap out the current symbolic state and cache it, like a swapfile for virtual memory
  - Caching formulas prevents re-execution, ensuring forward progress
- Concretization constraints prevent false positives
- Pure symbolic execution is hard
  - For example, memcmp creates a new branch for each character in the input buffer
  - No need to do symbolic execution for the setup before we are opening the program to any attacker input



# Bug Example

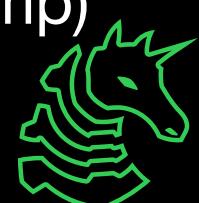
```
#define BUFSIZE 4096
typedef struct {
    char buf[BUFSIZE];
    int used;
} STATIC_BUFFER_t;
typedef struct conn {
    STATIC_BUFFER_t read_buf;
    ... // irrelevant
} CONN_t;
static void serverlog(LOG_TYPE_t type, const char* format, ...)
{
    ... // irrelevant
    if (format) {
        va_start (ap, format);
        vsprintf(buf, format, ap);
        va_end(ap);
    }
    fprintf(log, buf); // format string bug
    fflush(log);
}
```

```
HTTP_STATE_t http_read_request(CONN_t *conn) {
    ... // irrelevant
    while (conn->read_buf.used < BUFSIZE) {
        sz = static_buffer_read(conn, &conn->read_buf);
        if (sz < 0) {
            ...
            conn->read_buf.used += sz;
        if (memcmp(&conn->read_buf.buf[conn->read_buf.used]
                  - 4, "\r\n\r\n", 4) == 0 { break; }
        }
        if (conn->read_buf.used >= BUFSIZE) {
            conn->status.st = HTTP_STATUS_400;
            return HTTP_STATE_ERROR;
        }
        serverlog(ERROR_LOG, "%s\n", conn->read_buf.buf);
    } // serverlog user input comes from the outside (HTTP)
```



# How it works

- There's a concrete runner (client) and symbolic executor (server)
- Client has a taint tracker to see if user input does bad things
  - If so, pass it to the symbolic executor, and see if we can solve to do even worse things
  - Works by checking if any basic blocks are corrupted
- Server decides tasking for the client and where to explore
- So, taint tracker would see that we are calling `fprintf` with a user-specified string, which lets us control how we access stack memory
  - At this point, we would suspend concrete execution, and determine where we can branch to
  - Pathing to new places is modeled as additional constraints (one per jump)



# How it works

- When we hit any important parts of the program, the SES will try to build a solution to its symbolic equation
  - If it can find one, then any solution is, by necessity, a working exploit
  - It will crawl through each of these possibilities by context switching back to the client
- MAYHEM will continue to explore until either a solution is returned, we hit a user-specified maximum runtime, or all execution paths have been exhausted



# Results

- 29 exploitable vulnerabilities (2 0-days)
- Relatively limited bug classes (buffer overflow, format string, function pointer overwrite)
- Won in the 2016 DARPA Cyber Grand Challenge, although those binaries were relatively simplified compared to hardened real-world targets
- Darpa had the AlxCC challenge this year, which was like the 2016 Cyber Grand Challenge but using LLMs on top of these techniques



# Discussion

- Do you think symbolic execution can handle finding more complex bug classes, like heap exploits and logic bugs? How would you go about modeling a heap allocator or important logical state?
- Are there other new techniques that would help amplify the effectiveness of fuzzing or symbolic execution, or otherwise win back some of the intuition and creativity of a human exploit developer? LLMs have been surprisingly middling thus far...
- Which do you think moves faster, automated bug finding, or exploitation mitigation techniques? How would a modern redesign of this concept fare against modern security mitigations?



# Next Meetings

**2025-11-06 • This Thursday**

- Game Hacking
- Go write some cheats and finally learn how Windows works

**2025-11-09 • Next Sunday**

- Something social
- Your helper staff is too socially awkward to know what to do

**2025-11-13 • Next Thursday**

- Rubber Ducky / Bad USB
- Turn physical access into RCE with this one simple trick

